

Laboratorio 8
Exploración y Uso Avanzado de Plataformas IA, Repositorios
Profesionales y Herramientas Globales para el desarrollo de IA
y de SW

Materia:

Profundización de inteligencia artificial

Participantes:

Ana Maria Navarro

Profesor:

Carlos Betancourt correa

Universidad de Manizales
ingeniería en sistemas y telecomunicaciones
Manizales, Caldas, Colombia

2.6 Papers With Code

Laboratorio 8

The screenshot shows the Hugging Face platform interface. At the top, there's a search bar with the placeholder "Search models, datasets, users...". Below the search bar, a banner says "Hugging Face is way more fun with friends and colleagues! Join an organization". To the right of the banner are links for "Models", "Datasets", "Spaces", "Community", "Docs", "Enterprise", "Pricing", and a "Dismiss this message" button. A "Trending Papers" section is visible, with a search bar set to "text classification". Below the search bar, there are three paper cards:

- Arabic Synonym BERT-based Adversarial Examples for Text Classification** (3 authors, Feb 5, 2024)
- Self-Evolution Learning for Mixup: Enhance Data Augmentation on Few-Shot Text Classification Tasks** (7 authors, May 22, 2023)
- CoCoP: Enhancing Text Classification with LLM through Code Completion Prompt** (3 authors, Nov 13, 2024)

1. Introducción

El objetivo de esta actividad fue replicar parcialmente un experimento proveniente de la plataforma

Papers With Code...

Arabic Synonym BERT-based Adversarial Examples for Text Classification

Published on Feb 5, 2024

Authors: Norah Alshahrani, Saeid Alshahrani, Esma Wali, Jeanna Matthews

Abstract

A word-level study of adversarial attacks in Arabic using synonym-based adversarial examples suggests fine-tuned BERT models are more susceptible to these attacks than other models, though adversarial training can mitigate this vulnerability.

AI-generated summary

Text classification systems have been proven vulnerable to adversarial text examples, modified versions of the original text examples that are often unnoticed by human eyes, yet can force text classification models to alter their classification. Often, research works quantifying the impact of adversarial text attacks have been applied only to models trained in English. In this paper, we introduce the first word-level study of adversarial attacks in Arabic. Specifically, we use a synonym (word-level) attack using a Masked Language Modeling (MLM) task with a BERT model in a black-box setting to assess the robustness of the state-of-the-art text classification models to adversarial attacks in Arabic. To evaluate the grammatical and semantic similarities of the newly produced adversarial examples using our synonym BERT-based attack, we invite four human evaluators to assess and compare the produced adversarial examples with their original examples. We also study the transferability of these newly produced Arabic adversarial examples to various models and investigate the effectiveness of defense mechanisms against these adversarial examples on the BERT models. We find that fine-tuned BERT models were more susceptible to our synonym attacks than the other Deep Neural Networks (DNN) models like WordCNN and WordLSTM we trained. We also find that fine-tuned BERT models were more susceptible to transferred attacks. We, lastly, find that fine-tuned BERT models successfully regain at least 2% in accuracy after applying adversarial training as an initial defense mechanism.

[View arXiv page](#) [View PDF](#) [GitHub ★1](#) [Add to collection](#)

2. Enlace al paper y al repositorio

Paper: <https://paperswithcode.com>

Repositorio del código: GitHub del estudio.

3. Configuración del entorno

La replicación se realizó en Google Colab, configuración del entorno GPU y librerías necesarias.

4. Ejecución del Notebook

Se ejecutaron correctamente las primeras partes del notebook. Se obtuvo error en la falta de archivos CSV.

5. Análisis del Error Encontrado

Los archivos CSV no están incluidos en el repositorio...

6. Solución aplicada

Laboratorio 8

Se crearon archivos vacíos y se documentó el error como parte natural del proceso científico.

7. Comparación de Resultados (Paper vs Replicación)

Los resultados no pudieron obtenerse debido a la falta de los artefactos generados por los autores.

8. Conclusiones

La replicación parcial fue exitosa en reproducir el entorno, pero incompleta en resultados finales.

The screenshot shows a Jupyter Notebook interface with the following details:

- File Explorer:** Shows a directory structure with files like README.md, Transferability.ipynb, bert_synonym_attack.py, charsets.py, download_models.py, emojis.py, and word_tokenizer.py.
- Notebook View:** The notebook title is "bert_synonym_attack / Transferability.ipynb". It shows a preview of the code and a raw code editor.
- Code Content:** The code in the notebook is identical to the one shown in the file explorer, demonstrating the transferability attack.
- Terminal View:** The terminal tab shows the command "Transferability.ipynb" was run.
- Output View:** The output tab shows the execution results of the notebook code.

```
1 1
2 2
3 3
4 4
5 5
6 6
7 7
8 8
9 9
10 10
11 11
12 12
13 13
14 14
15 15
16 16
17 17
18 18
19 19
20 20
21 21
22 22
23 23
24 24
25 25
```

```
import os, time, torch
from transformers import logging as hf1
os.environ['TF_CPP_MIN_LOG_LEVEL'] = '2'
import logging, warnings, tensorflow as tf
from tensorflow_hub import disable_progress_bars
from tensorflow import enable_eager_execution
enable_eager_execution()

# disable_progress_bars()
start_time = time.time()
hf1.set_verbosity_error()
tf.autograph.set_verbosity(0)
tf.get_logger().setLevel('INFO')
tf.get_logger().setLevel(logging.ERROR)
os.environ['TOKENIZERS_PARALLELISM'] = 'False'
warnings.simplefilter(action='ignore', category=Warning)
warnings.simplefilter(action='ignore', category=FutureWarning)
tf.compat.v1.logging.set_verbosity(tf.compat.v1.logging.ERROR)

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

```
def load(model, dataset):
    if model == 'bert' and dataset == 'hard':
        from transformers import BertTokenizer, BertForSequenceClassification
        tokenizer = BertTokenizer.from_pretrained('models/bert-base-arabertv2')
        model = BertForSequenceClassification.from_pretrained('models/BERTHard', num_labels=4)
        return model, tokenizer

    elif model == 'cnn' and dataset == 'hard':
        import pickle
```



UNIVERSIDAD DE
MANIZALES®

Acreditación Institucional
de Alta Calidad
Resolución 4702 del 18 de mayo de 2019

Laboratorio 8

The screenshot shows a Jupyter Notebook interface with two code cells. The first cell, labeled [1], contains Python code related to setting up a torch environment and importing necessary modules. The second cell, labeled [2], contains a function definition for loading a model and dataset. A red box highlights the left margin of the first cell, where the output status '37 s' is displayed. A play button icon is visible at the bottom of the second cell's input area.

```
[1] ✓ 37 s
import os, time, torch
from transformers import logging as hf1
os.environ['TF_CPP_MIN_LOG_LEVEL'] = '2'
import logging, warnings, tensorflow as tf
from huggingface_hub.utils import disable_progress_bars

disable_progress_bars()
start_time = time.time()
hf1.set_verbosity_error()
tf.autograph.set_verbosity(0)
tf.get_logger().setLevel('INFO')
tf.get_logger().setLevel(logging.ERROR)
os.environ['TOKENIZERS_PARALLELISM'] = 'False'
warnings.simplefilter(action='ignore', category=Warning)
warnings.simplefilter(action='ignore', category=FutureWarning)
tf.compat.v1.logging.set_verbosity(tf.compat.v1.logging.ERROR)

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

[2] ✓ 0 s
```

```
def load(model, dataset):
    if model == 'bert' and dataset == 'hard':
        from transformers import BertTokenizer, BertForSequenceClassification
        tokenizer = BertTokenizer.from_pretrained('models/bert-base-arabertv2')
        model = BertForSequenceClassification.from_pretrained('models/BERThard', num_labels=4)
        return model, tokenizer

    elif model == 'cnn' and dataset == 'hard':
        import pickle
```